

```
In [1]: # 1. Load and Understand the Dataset
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.datasets import fetch_california_housing

data = fetch_california_housing(as_frame=True)
df = data.frame

print(df.head())
print("Shape:", df.shape)
print("Data types:\n", df.dtypes)
print("Missing values:\n", df.isnull().sum())

MedInc    HouseAge    AveRooms    AveBedrms    Population    AveOccup    Latitude \
0   8.3252      41.0   6.984127   1.023810      322.0   2.555556    37.88
1   8.3014      21.0   6.238137   0.971880     2401.0   2.109842    37.86
2   7.2574      52.0   8.288136   1.073446      496.0   2.802260    37.85
3   5.6431      52.0   5.817352   1.073059      558.0   2.547945    37.85
4   3.8462      52.0   6.281853   1.081081      565.0   2.181467    37.85

Longitude    MedHouseVal
0   -122.23        4.526
1   -122.22        3.585
2   -122.24        3.521
3   -122.25        3.413
4   -122.25        3.422
Shape: (20640, 9)
Data types:
MedInc           float64
HouseAge          float64
AveRooms          float64
AveBedrms          float64
Population         float64
AveOccup          float64
Latitude           float64
Longitude          float64
MedHouseVal        float64
dtype: object
Missing values:
MedInc            0
HouseAge           0
AveRooms           0
AveBedrms          0
Population          0
AveOccup           0
Latitude            0
Longitude           0
MedHouseVal         0
dtype: int64
```

```
In [2]: # 2. Exploratory Data Analysis (EDA)
plt.figure(figsize=(6,4))
sns.histplot(df['MedHouseVal'], bins=30, kde=True)
plt.title("Distribution of Median House Value")
plt.show()


# Scatter plot:
```

```
sns.scatterplot(x='MedInc', y='MedHouseVal', data=df)
plt.title("Median Income vs House Value")
plt.show()


# Correlation heatmap:
plt.figure(figsize=(10,6))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
plt.title("Correlation Heatmap")
plt.show()

# Boxplot:
df['IncomeBin'] = pd.qcut(df['MedInc'], q=4)
plt.figure(figsize=(8,5))
sns.boxplot(x='IncomeBin', y='MedHouseVal', data=df)
plt.title("House Value by Income Quartile")
plt.xticks(rotation=45)
plt.show()


# Visual outlier detection
plt.figure(figsize=(6,4))
sns.boxplot(x=df['MedHouseVal'])
plt.title("Boxplot of Median House Value")
plt.show()
```




Median Income vs House Value




Correlation Heatmap



House Value by Income Quartile



Boxplot of Median House Value



```
In [9]: # 3. Data Preprocessing
```

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

# Select features
features = ['MedInc', 'AveRooms', 'AveOccup', 'HouseAge', 'AveBedrms']
X = df[features]
```

```

y = df['MedHouseVal']

# Scaling
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2,
print(X_train, X_test, y_train, y_test)
[[ -0.32165429 -0.16625886  0.05980847  0.34647803 -0.19045099]
 [-0.03061993 -0.3861807 -0.12830597  1.61780729 -0.1174719 ]
 [ 0.1503488  0.08764093 -0.03345346 -1.95780625 -0.2354004 ]
 ...
 [-0.49283239 -0.58294927  0.02517025  0.58485227 -0.03582776]
 [ 0.97302487  0.39058403  0.01042151 -1.08376738 -0.06055424]
 [-0.68174943 -0.81905034 -0.09262259  1.85618152 -0.07997264]] [[-1.15248922 -0.
49989596  0.07768129 -0.2891866 -0.15697642]
 [-0.70501534 -0.1574771 -0.03763415  0.10810379  0.20430076]
 [-0.20558796 -0.5868143 -0.164679  1.85618152  0.1882305 ]
 ...
 [ 2.82092723  0.73104024 -0.02697551 -0.2891866 -0.31545104]
 [-0.57147326 -0.05657389 -0.04641411  0.58485227 -0.23961371]
 [-0.16768875 -0.58209158  0.06347664 -0.92485123 -0.13334689]] 14196      1.030
8267      3.821
17445     1.726
14265     0.934
2271      0.965
...
11284      2.292
11964     0.978
5390       2.221
860        2.835
15795     3.250
Name: MedHouseVal, Length: 16512, dtype: float64 20046      0.47700
3024      0.45800
15663     5.00001
20484     2.18600
9814      2.78000
...
15362     2.63300
16623     2.66800
18086     5.00001
2144      0.72300
3665      1.51500
Name: MedHouseVal, Length: 4128, dtype: float64

```

In [5]: # 4 & 5. Build and Train Two Regression Models

```

from sklearn.linear_model import LinearRegression

# Model A: Simple Linear Regression
X_simple = df[['MedInc']]
X_simple_scaled = scaler.fit_transform(X_simple)
X_train_s, X_test_s, y_train_s, y_test_s = train_test_split(X_simple_scaled, y,

model_a = LinearRegression()
model_a.fit(X_train_s, y_train_s)

# Model B: Multiple Linear Regression

```

```
model_b = LinearRegression()
model_b.fit(X_train, y_train)
```

Out[5]:

▼ LinearRegression ⓘ ⓘ

► Parameters

In [6]: # 6. Evaluate the Models

```
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

def evaluate(model, X_test, y_test, name="Model"):
    y_pred = model.predict(X_test)
    print(f"\n{name} Evaluation:")
    print("MAE:", mean_absolute_error(y_test, y_pred))
    print("MSE:", mean_squared_error(y_test, y_pred))
    print("RMSE:", np.sqrt(mean_squared_error(y_test, y_pred)))
    print("R2 Score:", r2_score(y_test, y_pred))
    return y_pred

y_pred_a = evaluate(model_a, X_test_s, y_test_s, "Simple Linear Regression")
y_pred_b = evaluate(model_b, X_test, y_test, "Multiple Linear Regression")
```

Simple Linear Regression Evaluation:

MAE: 0.629908653009376
 MSE: 0.7091157771765548
 RMSE: 0.8420901241414454
 R2 Score: 0.45885918903846656

Multiple Linear Regression Evaluation:

MAE: 0.5795921844399792
 MSE: 0.6423670622489219
 RMSE: 0.8014780485134462
 R2 Score: 0.5097965040568945



In [7]: # 8. Required Visualizations

```
# Scatter plot + regression line for simple regression
plt.figure(figsize=(6,4))
sns.regplot(x=X_test_s.flatten(), y=y_test_s, line_kws={"color": "red"})
plt.title("Simple Regression: MedInc vs House Value")
plt.xlabel("Median Income (scaled)")
plt.ylabel("House Value")
plt.show()


# Actual vs Predicted plot for multiple regression
plt.figure(figsize=(6,4))
sns.scatterplot(x=y_test, y=y_pred_b)
plt.plot([y.min(), y.max()], [y.min(), y.max()], 'r--')
plt.xlabel("Actual")
plt.ylabel("Predicted")
plt.title("Actual vs Predicted (Multiple Regression)")
plt.show()

# Residual plot
residuals = y_test - y_pred_b
plt.figure(figsize=(6,4))
sns.scatterplot(x=y_pred_b, y=residuals)
plt.axhline(0, color='red', linestyle='--')
```

```
plt.title("Residuals vs Predicted")
plt.xlabel("Predicted")
plt.ylabel("Residuals")
plt.show()
```



Residuals vs Predicted



In []: